

# MFCC 10 sec Model Creation

May 22, 2023

```
[1]: import os
import json
import csv
import sys
import numpy as np
import matplotlib.pyplot as plt
import librosa
import tensorflow as tf
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit
from tensorflow.keras.utils import plot_model
import datetime
import time
from tensorflow.keras.callbacks import TensorBoard
```

2023-05-16 22:04:42.397162: I tensorflow/core/platform/cpu\_feature\_guard.cc:182]  
This TensorFlow binary is optimized to use available CPU instructions in  
performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild  
TensorFlow with the appropriate compiler flags.

```
[2]: csv_path = "csv_fma_10secs_data.csv"
sr=22050
csv.field_size_limit(sys.maxsize)
```

[2]: 131072

```
[3]: #Read MFCC data from CSV
def csv_read_data(csv_path):
    # Load data from CSV file
    with open(csv_path, 'r') as csvfile:
        reader = csv.reader(csvfile)

        # Skip the header row
        next(reader)

        # Initialize lists to hold genre and MFCC data
        genres = []
        mfcc = []
```

```

    # Iterate over each row of the CSV file
    for row in reader:
        # Extract genre and MFCC data from the row
        genre = int(row[0])
        mfcc_data = json.loads(row[1])

        # Append genre and MFCC data to lists
        genres.append(genre)
        mfcc.append(mfcc_data)

    # Convert lists to numpy arrays
    X = np.array(mfcc)
    y = np.array(genres)

    return X, y

```

```

[4]: def stratified_split_dataset(inputs, targets, test_split_size, val_split_size):
    sss = StratifiedShuffleSplit(n_splits=1, test_size=test_split_size,
    ↪random_state=42)

    train_val_indices, test_indices = next(sss.split(inputs, targets))
    inputs_train_val, targets_train_val = inputs[train_val_indices],
    ↪targets[train_val_indices]

    val_size = val_split_size / (1 - test_split_size)
    sss_train = StratifiedShuffleSplit(n_splits=1, test_size=val_size,
    ↪random_state=43)
    train_indices, val_indices = next(sss_train.split(inputs_train_val,
    ↪targets_train_val))

    inputs_train, targets_train = inputs_train_val[train_indices],
    ↪targets_train_val[train_indices]
    inputs_val, targets_val = inputs_train_val[val_indices],
    ↪targets_train_val[val_indices]
    inputs_test, targets_test = inputs[test_indices], targets[test_indices]

    # Needed for compatibility reasons
    inputs_train = inputs_train[..., np.newaxis]
    inputs_val = inputs_val[..., np.newaxis]
    inputs_test = inputs_test[..., np.newaxis]

    return inputs_train, inputs_val, inputs_test, targets_train, targets_val,
    ↪targets_test

```

[5]: *#Used to plot performance history of model*

```
def plot_performance(hist):  
  
    acc = hist.history['acc']  
    val_acc = hist.history['val_acc']  
    loss = hist.history['loss']  
    val_loss = hist.history['val_loss']  
  
    epochs = range(len(acc))  
  
    plt.plot(epochs, acc, 'r', label='Training accuracy')  
    plt.plot(epochs, val_acc, 'b', label='Validation accuracy')  
    plt.title('Training and validation accuracy')  
    plt.legend()  
    plt.figure()  
  
    plt.plot(epochs, loss, 'r', label='Training Loss')  
    plt.plot(epochs, val_loss, 'b', label='Validation Loss')  
    plt.title('Training and validation loss')  
    plt.legend()  
  
    plt.show()
```

[6]: *#===== MAIN =====*

```
#Reading MFCC data from CSV  
inputs, targets = csv_read_data(csv_path)
```

[7]: *#Splitting data*

```
Xtrain, Xval, Xtest, ytrain, yval, ytest =stratified_split_dataset(inputs,   
    ↪targets, 0.15, 0.15)
```

[8]: 

```
print(Xtrain.shape)  
print(Xtest.shape)  
print(Xval.shape)
```

```
(7301, 417, 40, 1)  
(1565, 417, 40, 1)  
(1565, 417, 40, 1)
```

[9]: *#Design the model*

```
def design_model(input_shape):  
  
    model = tf.keras.models.Sequential([  
  
        tf.keras.layers.Conv2D(64, (3,3), activation='relu',   
    ↪input_shape=input_shape),  
        tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
```

```

tf.keras.layers.BatchNormalization(),

tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
tf.keras.layers.BatchNormalization(),

tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
tf.keras.layers.BatchNormalization(),
tf.keras.layers.Dropout(0.3),

tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
tf.keras.layers.Dropout(0.3),

tf.keras.layers.Flatten(),
tf.keras.layers.Dense(64, activation='relu'),
tf.keras.layers.Dense(len(np.unique(targets)), activation='softmax')
])

return model

```

```

[10]: #Model Creation
input_shape = (Xtrain.shape[1], Xtrain.shape[2], 1)
model = design_model(input_shape)

model.compile(optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001),
              loss = 'sparse_categorical_crossentropy',
              metrics = ['acc'])

```

```

2023-05-16 22:06:34.848421: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1635] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 22077 MB memory: -> device:
0, name: NVIDIA GeForce RTX 4090, pci bus id: 0000:c1:00.0, compute capability:
8.9

```

```

[11]: #Model Summary
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 415, 38, 64)	640
max_pooling2d (MaxPooling2D)	(None, 208, 19, 64)	0

batch_normalization (Batch Normalization)	(None, 208, 19, 64)	256
conv2d_1 (Conv2D)	(None, 206, 17, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 103, 9, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 103, 9, 64)	256
conv2d_2 (Conv2D)	(None, 101, 7, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(None, 51, 4, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 51, 4, 64)	256
dropout (Dropout)	(None, 51, 4, 64)	0
conv2d_3 (Conv2D)	(None, 49, 2, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 25, 1, 64)	0
dropout_1 (Dropout)	(None, 25, 1, 64)	0
flatten (Flatten)	(None, 1600)	0
dense (Dense)	(None, 64)	102464
dense_1 (Dense)	(None, 6)	390

```
=====
Total params: 215,046
Trainable params: 214,662
Non-trainable params: 384
-----
```

```
[12]: #TENSORBOARD
NAME= "MFCC 10secs Model"
log_dir = "logs/fit/" + NAME
tensorboard_callback = [tf.keras.callbacks.TensorBoard(log_dir=log_dir,
↪ histogram_freq=1),tf.keras.callbacks.EarlyStopping(patience=5,
↪ restore_best_weights=True)]
```

```
[13]: #COMPILATION
history = model.fit(Xtrain, ytrain,
                    validation_data = (Xval, yval),
                    epochs = 550,
                    batch_size = 64,
                    callbacks=[tensorboard_callback])

model.save("./Models/"+NAME)
```

Epoch 1/550

```
2023-05-16 22:06:36.657097: E
tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] layout failed:
INVALID_ARGUMENT: Size of values 0 does not match size of permutation 4 @ fanin
shape insequential/dropout/dropout/SelectV2-2-TransposeNHWCToNCHW-
LayoutOptimizer
2023-05-16 22:06:37.598943: I
tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:424] Loaded cuDNN
version 8600
2023-05-16 22:06:38.333676: I
tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:637] TensorFloat-32
will be used for the matrix multiplication. This will only be logged once.
2023-05-16 22:06:38.505280: I tensorflow/compiler/xla/service/service.cc:169]
XLA service 0x7f2301ddee10 initialized for platform CUDA (this does not
guarantee that XLA will be used). Devices:
2023-05-16 22:06:38.505309: I tensorflow/compiler/xla/service/service.cc:177]
StreamExecutor device (0): NVIDIA GeForce RTX 4090, Compute Capability 8.9
2023-05-16 22:06:38.509082: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR
crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
2023-05-16 22:06:38.637358: I ./tensorflow/compiler/jit/device_compiler.h:180]
Compiled cluster using XLA! This line is logged at most once for the lifetime
of the process.
```

```
115/115 [=====] - 7s 17ms/step - loss: 1.6261 - acc:
0.4087 - val_loss: 1.4462 - val_acc: 0.4102
```

Epoch 2/550

```
115/115 [=====] - 1s 12ms/step - loss: 1.2304 - acc:
0.5388 - val_loss: 1.1591 - val_acc: 0.5917
```

Epoch 3/550

```
115/115 [=====] - 1s 12ms/step - loss: 1.1372 - acc:
0.5742 - val_loss: 1.0622 - val_acc: 0.6083
```

Epoch 4/550

```
115/115 [=====] - 1s 13ms/step - loss: 1.0800 - acc:
0.5940 - val_loss: 1.0237 - val_acc: 0.6096
```

Epoch 5/550

```
115/115 [=====] - 1s 13ms/step - loss: 1.0405 - acc:
0.6143 - val_loss: 1.0031 - val_acc: 0.6339
```

Epoch 6/550  
115/115 [=====] - 1s 12ms/step - loss: 1.0134 - acc:  
0.6191 - val\_loss: 0.9981 - val\_acc: 0.6166  
Epoch 7/550  
115/115 [=====] - 1s 13ms/step - loss: 0.9813 - acc:  
0.6357 - val\_loss: 0.9589 - val\_acc: 0.6441  
Epoch 8/550  
115/115 [=====] - 1s 13ms/step - loss: 0.9629 - acc:  
0.6410 - val\_loss: 1.0149 - val\_acc: 0.6224  
Epoch 9/550  
115/115 [=====] - 1s 12ms/step - loss: 0.9514 - acc:  
0.6480 - val\_loss: 0.9490 - val\_acc: 0.6543  
Epoch 10/550  
115/115 [=====] - 1s 12ms/step - loss: 0.9259 - acc:  
0.6551 - val\_loss: 0.9286 - val\_acc: 0.6594  
Epoch 11/550  
115/115 [=====] - 1s 12ms/step - loss: 0.9031 - acc:  
0.6651 - val\_loss: 0.9384 - val\_acc: 0.6498  
Epoch 12/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8888 - acc:  
0.6705 - val\_loss: 0.9153 - val\_acc: 0.6588  
Epoch 13/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8655 - acc:  
0.6737 - val\_loss: 0.9234 - val\_acc: 0.6537  
Epoch 14/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8500 - acc:  
0.6805 - val\_loss: 0.9081 - val\_acc: 0.6690  
Epoch 15/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8484 - acc:  
0.6852 - val\_loss: 0.9301 - val\_acc: 0.6486  
Epoch 16/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8148 - acc:  
0.6957 - val\_loss: 0.9086 - val\_acc: 0.6703  
Epoch 17/550  
115/115 [=====] - 1s 12ms/step - loss: 0.8062 - acc:  
0.6978 - val\_loss: 0.8751 - val\_acc: 0.6882  
Epoch 18/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7860 - acc:  
0.7002 - val\_loss: 0.8630 - val\_acc: 0.6920  
Epoch 19/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7743 - acc:  
0.7120 - val\_loss: 0.8753 - val\_acc: 0.6843  
Epoch 20/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7559 - acc:  
0.7163 - val\_loss: 0.8842 - val\_acc: 0.6760  
Epoch 21/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7530 - acc:  
0.7192 - val\_loss: 0.8619 - val\_acc: 0.6741

Epoch 22/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7350 - acc:  
0.7169 - val\_loss: 0.8635 - val\_acc: 0.6927

Epoch 23/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7225 - acc:  
0.7262 - val\_loss: 0.8823 - val\_acc: 0.6786

Epoch 24/550  
115/115 [=====] - 1s 12ms/step - loss: 0.7023 - acc:  
0.7404 - val\_loss: 0.8693 - val\_acc: 0.6818

Epoch 25/550  
115/115 [=====] - 1s 12ms/step - loss: 0.6879 - acc:  
0.7406 - val\_loss: 0.8346 - val\_acc: 0.7093

Epoch 26/550  
115/115 [=====] - 1s 12ms/step - loss: 0.6732 - acc:  
0.7462 - val\_loss: 0.8621 - val\_acc: 0.6767

Epoch 27/550  
115/115 [=====] - 1s 12ms/step - loss: 0.6583 - acc:  
0.7562 - val\_loss: 0.8584 - val\_acc: 0.6888

Epoch 28/550  
115/115 [=====] - 1s 12ms/step - loss: 0.6472 - acc:  
0.7592 - val\_loss: 0.8233 - val\_acc: 0.7042

Epoch 29/550  
115/115 [=====] - 1s 12ms/step - loss: 0.6269 - acc:  
0.7678 - val\_loss: 0.8469 - val\_acc: 0.7029

Epoch 30/550  
115/115 [=====] - 1s 13ms/step - loss: 0.6141 - acc:  
0.7711 - val\_loss: 0.8325 - val\_acc: 0.6965

Epoch 31/550  
115/115 [=====] - 1s 13ms/step - loss: 0.6069 - acc:  
0.7754 - val\_loss: 0.8664 - val\_acc: 0.6831

Epoch 32/550  
115/115 [=====] - 1s 13ms/step - loss: 0.6033 - acc:  
0.7741 - val\_loss: 0.8343 - val\_acc: 0.6952

Epoch 33/550  
115/115 [=====] - 1s 13ms/step - loss: 0.5942 - acc:  
0.7792 - val\_loss: 0.8157 - val\_acc: 0.7093

Epoch 34/550  
115/115 [=====] - 1s 13ms/step - loss: 0.5615 - acc:  
0.7911 - val\_loss: 0.8282 - val\_acc: 0.7042

Epoch 35/550  
115/115 [=====] - 1s 13ms/step - loss: 0.5501 - acc:  
0.7930 - val\_loss: 0.8143 - val\_acc: 0.7022

Epoch 36/550  
115/115 [=====] - 1s 13ms/step - loss: 0.5466 - acc:  
0.8000 - val\_loss: 0.8151 - val\_acc: 0.7054

Epoch 37/550  
115/115 [=====] - 2s 13ms/step - loss: 0.5247 - acc:  
0.8076 - val\_loss: 0.8628 - val\_acc: 0.6927



```

Epoch 38/550
115/115 [=====] - 1s 13ms/step - loss: 0.5188 - acc:
0.8062 - val_loss: 0.8348 - val_acc: 0.7048
Epoch 39/550
115/115 [=====] - 1s 12ms/step - loss: 0.5145 - acc:
0.8082 - val_loss: 0.8309 - val_acc: 0.7042
Epoch 40/550
115/115 [=====] - 1s 12ms/step - loss: 0.5039 - acc:
0.8163 - val_loss: 0.8060 - val_acc: 0.7137
Epoch 41/550
115/115 [=====] - 1s 12ms/step - loss: 0.4769 - acc:
0.8233 - val_loss: 0.8245 - val_acc: 0.7029
Epoch 42/550
115/115 [=====] - 1s 13ms/step - loss: 0.4886 - acc:
0.8132 - val_loss: 0.8288 - val_acc: 0.7073
Epoch 43/550
115/115 [=====] - 1s 13ms/step - loss: 0.4511 - acc:
0.8358 - val_loss: 0.8501 - val_acc: 0.7010
Epoch 44/550
115/115 [=====] - 1s 13ms/step - loss: 0.4546 - acc:
0.8380 - val_loss: 0.8297 - val_acc: 0.7054
Epoch 45/550
115/115 [=====] - 1s 12ms/step - loss: 0.4373 - acc:
0.8328 - val_loss: 0.8500 - val_acc: 0.7105

2023-05-16 22:07:45.732154: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,51,4,64]
[[{{node inputs}}]]
2023-05-16 22:07:45.745121: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,25,1,64]
[[{{node inputs}}]]
2023-05-16 22:07:46.206991: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,51,4,64]
[[{{node inputs}}]]
2023-05-16 22:07:46.240386: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,25,1,64]
[[{{node inputs}}]]
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op,
_jit_compiled_convolution_op while saving (showing 4 of 4). These functions will


```


not be directly callable after loading.

INFO:tensorflow:Assets written to: ./Models/MFCC 10secs Model/assets

INFO:tensorflow:Assets written to: ./Models/MFCC 10secs Model/assets

```
[14]: #Plot Performance  
plot_performance(history)
```





```
[15]: y_pred = model.predict(Xtest)
```

49/49 [=====] - 0s 2ms/step

```
[16]: from sklearn.metrics import accuracy_score

y_pred_classes = np.argmax(y_pred, axis=1)

# Calculate the accuracy
accuracy = accuracy_score(ytest, y_pred_classes)
print(f"Test accuracy: {accuracy:.4f}")
```

Test accuracy: 0.7016

```
[19]: from sklearn.metrics import f1_score
```

```
[22]: # Make predictions on the test set
y_pred_test = model.predict(Xtest)

# Convert probabilities to class labels
y_pred_test_labels = np.argmax(y_pred_test, axis=1)
```

```
# Compute the F1 score
f1_test = f1_score(ytest, y_pred_test_labels, average='micro')

print('Test F1 score: ', f1_test)
```

49/49 [=====] - 0s 2ms/step  
Test F1 score: 0.7015974440894569