

MFCC Model Final

May 6, 2023

```
[1]: import os
import json
import csv
import sys
import numpy as np
import matplotlib.pyplot as plt
import librosa
import tensorflow as tf
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit
from tensorflow.keras.utils import plot_model
import datetime
import time
from tensorflow.keras.callbacks import TensorBoard
```

2023-05-04 12:10:32.231539: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.

```
[2]: #Paths and declaration
audio_path = "./FMA Balanced"
csv_path = "csv_data.csv"

sr=22050

csv.field_size_limit(sys.maxsize)
```

[2]: 131072

```
[3]: #All songs under 30 seconds
TOTAL_SAMPLES = 29 * sr

#Split the song in 3
NUM_SLICES = 3
SAMPLES_PER_SLICE = int(TOTAL_SAMPLES / NUM_SLICES)
```

```
[4]: #Write MFCC to CSV
def csv_write_data(audio_path, csv_path):
    #Create a list of rows for writing to CSV
    rows = []

    #Generate MFCC for every song & writing to CSV
    for i, (dirpath, dirnames, filenames) in enumerate(os.walk(audio_path)):
        for file in filenames:
            #Generate MFCC for every item in audio_path * slice
            song, sr = librosa.load(os.path.join(dirpath, file), duration=29)

            for s in range(NUM_SLICES):
                start_sample = SAMPLES_PER_SLICE * s
                end_sample = start_sample + SAMPLES_PER_SLICE

                mfcc = librosa.feature.mfcc(y=song[start_sample:end_sample],
↪sr=sr, n_mfcc=40)
                mfcc = mfcc.T

                rows.append([i-1, json.dumps(mfcc.tolist())])

    #Writing CSV
    with open(csv_path, 'w', newline='') as csvfile:
        writer = csv.writer(csvfile)

        # Write the header row
        writer.writerow(["genre", "mfcc"])

        # Write the values rows
        for row in rows:
            writer.writerow(row)
```

```
[5]: #Read MFCC data from CSV
def csv_read_data(csv_path):
    # Load data from CSV file
    with open(csv_path, 'r') as csvfile:
        reader = csv.reader(csvfile)

        next(reader)

        # Initialize lists to hold genre and MFCC data
        genres = []
        mfcc = []

        # Iterate over each row of the CSV file
        for row in reader:
```

```

        # Extract genre and MFCC data from the row
        genre = int(row[0])
        mfcc_data = json.loads(row[1])

        # Append genre and MFCC data to lists
        genres.append(genre)
        mfcc.append(mfcc_data)

    # Convert lists to numpy arrays
    X = np.array(mfcc)
    y = np.array(genres)

    return X, y

```

[6]: *#Use stratified split to separate data into training, validation and test*

```

def stratified_split_dataset(inputs, targets, split_size):
    sss = StratifiedShuffleSplit(n_splits=1, test_size=split_size,
    ↪random_state=42)

    train_val_indices, test_indices = next(sss.split(inputs, targets))
    inputs_train_val, targets_train_val = inputs[train_val_indices],
    ↪targets[train_val_indices]

    sss_train = StratifiedShuffleSplit(n_splits=1, test_size=split_size,
    ↪random_state=43)
    train_indices, val_indices = next(sss_train.split(inputs_train_val,
    ↪targets_train_val))

    inputs_train, targets_train = inputs_train_val[train_indices],
    ↪targets_train_val[train_indices]
    inputs_val, targets_val = inputs_train_val[val_indices],
    ↪targets_train_val[val_indices]
    inputs_test, targets_test = inputs[test_indices], targets[test_indices]

    #Needed for compatibility reasons
    inputs_train = inputs_train[...,np.newaxis]
    inputs_val = inputs_val[...,np.newaxis]
    inputs_test= inputs_test[..., np.newaxis]

    return inputs_train, inputs_val, inputs_test, targets_train, targets_val,
    ↪targets_test

```

[15]: *#Design the model*

```

def design_model(input_shape):

```

```

model = tf.keras.models.Sequential([

    tf.keras.layers.Conv2D(64, (3,3), activation='relu',
↪input_shape=input_shape),
    tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
    tf.keras.layers.BatchNormalization(),

    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
    tf.keras.layers.BatchNormalization(),

    tf.keras.layers.Conv2D(64, (2,2), activation='relu'),
    tf.keras.layers.MaxPooling2D((3,3), strides=(2,2), padding='same'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.3),

    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(len(np.unique(targets)), activation='softmax')
])

return model

```

```

[8]: #Used to plot performance history of model
def plot_performance(hist):

    acc = hist.history['acc']
    val_acc = hist.history['val_acc']
    loss = hist.history['loss']
    val_loss = hist.history['val_loss']

    epochs = range(len(acc))

    plt.plot(epochs, acc, 'r', label='Training accuracy')
    plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
    plt.title('Training and validation accuracy')
    plt.legend()
    plt.figure()

    plt.plot(epochs, loss, 'r', label='Training Loss')
    plt.plot(epochs, val_loss, 'b', label='Validation Loss')
    plt.title('Training and validation loss')
    plt.legend()

    plt.show()

```

```
[9]: #===== MAIN =====
```

```
#CSV WRITING DATA
```

```
csv_write_data(audio_path, csv_path)
```

```
[10]: #Reading MFCC data from CSV
```

```
inputs, targets = csv_read_data(csv_path)
```

```
[11]: #Splitting data
```

```
Xtrain, Xval, Xtest, ytrain, yval, ytest =stratified_split_dataset(inputs, targets, 0.2)
```

```
[16]: #Model Creation
```

```
input_shape = (Xtrain.shape[1], Xtrain.shape[2], 1)
```

```
model = design_model(input_shape)
```

```
model.compile(optimizer = tf.keras.optimizers.Adam(learning_rate=0.001),  
              loss = 'sparse_categorical_crossentropy',  
              metrics = ['acc'])
```

2023-05-04 12:10:59.823788: I

tensorflow/core/common_runtime/gpu/gpu_device.cc:1635] Created device

/job:localhost/replica:0/task:0/device:GPU:0 with 22250 MB memory: -> device:
0, name: NVIDIA GeForce RTX 3090, pci bus id: 0000:81:00.0, compute capability:
8.6

```
[22]: #Model Summary
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 415, 38, 64)	640
max_pooling2d (MaxPooling2D)	(None, 208, 19, 64)	0
batch_normalization (Batch Normalization)	(None, 208, 19, 64)	256
conv2d_1 (Conv2D)	(None, 206, 17, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 103, 9, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 103, 9, 64)	256

```

hNormalization)

conv2d_2 (Conv2D)          (None, 102, 8, 64)      16448

max_pooling2d_2 (MaxPooling (None, 51, 4, 64)      0
2D)

batch_normalization_2 (Batc (None, 51, 4, 64)      256
hNormalization)

dropout (Dropout)          (None, 51, 4, 64)      0

flatten (Flatten)          (None, 13056)           0

dense (Dense)              (None, 64)             835648

dense_1 (Dense)            (None, 7)              455

```

```

=====
Total params: 890,887
Trainable params: 890,503
Non-trainable params: 384
-----

```

```

[17]: #TENSORBOARD
NAME= "Final Version"
log_dir = "logs/fit/" + NAME
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
↪ histogram_freq=1)

```

```

[18]: #COMPILATION
history = model.fit(Xtrain, ytrain,
                    validation_data = (Xval, yval),
                    epochs = 15,
                    batch_size = 64,
                    callbacks=[tensorboard_callback])

```

Epoch 1/15

```

2023-05-04 12:11:01.034813: E
tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] layout failed:
INVALID_ARGUMENT: Size of values 0 does not match size of permutation 4 @ fanin
shape insequential/dropout/dropout/SelectV2-2-TransposeNHWCToNCHW-
LayoutOptimizer
2023-05-04 12:11:02.000143: I
tensorflow/compiler/xla/stream_executor/cuda/cuda_dnn.cc:424] Loaded cuDNN
version 8600
2023-05-04 12:11:02.704844: I

```

```

tensorflow/compiler/xla/stream_executor/cuda/cuda_blas.cc:637] TensorFlow-32
will be used for the matrix multiplication. This will only be logged once.
2023-05-04 12:11:02.706143: I tensorflow/compiler/xla/service/service.cc:169]
XLA service 0x7fe0c6dc0f70 initialized for platform CUDA (this does not
guarantee that XLA will be used). Devices:
2023-05-04 12:11:02.706160: I tensorflow/compiler/xla/service/service.cc:177]
StreamExecutor device (0): NVIDIA GeForce RTX 3090, Compute Capability 8.6
2023-05-04 12:11:02.710435: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR
crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
2023-05-04 12:11:02.853797: I ./tensorflow/compiler/jit/device_compiler.h:180]
Compiled cluster using XLA! This line is logged at most once for the lifetime
of the process.

```

```

21/21 [=====] - 5s 35ms/step - loss: 1.9201 - acc:
0.3475 - val_loss: 4.7497 - val_acc: 0.2560

```

Epoch 2/15

```

21/21 [=====] - 0s 23ms/step - loss: 1.3210 - acc:
0.5045 - val_loss: 3.2636 - val_acc: 0.2827

```

Epoch 3/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.9974 - acc:
0.6153 - val_loss: 2.9248 - val_acc: 0.3185

```

Epoch 4/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.7149 - acc:
0.7396 - val_loss: 2.6424 - val_acc: 0.3542

```

Epoch 5/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.5168 - acc:
0.8304 - val_loss: 2.1787 - val_acc: 0.3452

```

Epoch 6/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.3119 - acc:
0.9174 - val_loss: 1.7571 - val_acc: 0.4613

```

Epoch 7/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.1776 - acc:
0.9591 - val_loss: 1.8806 - val_acc: 0.5089

```

Epoch 8/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.1116 - acc:
0.9814 - val_loss: 1.5899 - val_acc: 0.5476

```

Epoch 9/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.0679 - acc:
0.9918 - val_loss: 1.4615 - val_acc: 0.5774

```

Epoch 10/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.0392 - acc:
0.9970 - val_loss: 1.4996 - val_acc: 0.5714

```

Epoch 11/15

```

21/21 [=====] - 0s 22ms/step - loss: 0.0225 - acc:
0.9993 - val_loss: 1.4392 - val_acc: 0.5804

```

Epoch 12/15


```


21/21 [=====] - 0s 21ms/step - loss: 0.0156 - acc:

```

```
0.9993 - val_loss: 1.4398 - val_acc: 0.5923
Epoch 13/15
21/21 [=====] - 0s 21ms/step - loss: 0.0134 - acc:
1.0000 - val_loss: 1.4690 - val_acc: 0.5952
Epoch 14/15
21/21 [=====] - 0s 22ms/step - loss: 0.0093 - acc:
1.0000 - val_loss: 1.4324 - val_acc: 0.6250
Epoch 15/15
21/21 [=====] - 0s 22ms/step - loss: 0.0067 - acc:
1.0000 - val_loss: 1.4887 - val_acc: 0.6250
```

```
[19]: #Plot Performance
      plot_performance(history)
```





```
[20]: model.save(NAME)
```

```
2023-05-04 12:11:12.700613: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,51,4,64]
```

```
[[{{node inputs}}]]
```

```
2023-05-04 12:11:13.229436: I tensorflow/core/common_runtime/executor.cc:1197]
[/device:CPU:0] (DEBUG INFO) Executor start aborting (this does not indicate an
error and you can ignore this message): INVALID_ARGUMENT: You must feed a value
for placeholder tensor 'inputs' with dtype float and shape [?,51,4,64]
```

```
[[{{node inputs}}]]
```

```
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op,
_jit_compiled_convolution_op, _jit_compiled_convolution_op while saving (showing
3 of 3). These functions will not be directly callable after loading.
```

```
INFO:tensorflow:Assets written to: Random Sampler Re run4/assets
```

```
INFO:tensorflow:Assets written to: Random Sampler Re run4/assets
```

```
[21]: model.save(NAME+".h5", save_format='h5')
```